

Rice Crop Yield Prediction

The goal of this project is to find insights on predicting crop yields of rice in India. Smallholder farmers, who make up most of the workforce there, face challenges like poverty and malnutrition. Predictive insights can help them make better decisions about resource use, and help in supporting food security.

The **dataset** can be found on [Kaggle](#)

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, MinMaxScaler, OneHotEncoder,
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, roc_curve, auc, classification_r
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn import tree
from sklearn.model_selection import cross_val_score, cross_validate
import imblearn
from imblearn.over_sampling import SMOTE, SMOTENC
from imblearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.compose import ColumnTransformer, make_column_selector
from sklearn.impute import SimpleImputer
from sklearn.model_selection import GridSearchCV
from sklearn.tree import export_text, plot_tree
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.preprocessing import FunctionTransformer
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.compose import TransformedTargetRegressor
from sklearn.feature_selection import VarianceThreshold
from lime.lime_tabular import LimeTabularExplainer
```

Let's load and examine the data

```
In [2]: # to show all columns
pd.set_option('display.max_columns', None)
```

```
# to revert: pd.reset_option('display.max_columns')
```

```
In [3]: df = pd.read_csv("data/Train.csv", index_col=0)
df.head(12)
```

Out[3]:

	District	Block	CultLand	CropCultLand	LandPreparationMet
ID					
ID_GTFAC7PEVWQ9	Nalanda	Noorsarai	45	40	TractorPlo FourWheelTracRotav
ID_TK40ARLSPOKS	Nalanda	Rajgir	26	26	WetTillagePudc TractorPlo FourWheelTr
ID_1FJY2CRIMLZZ	Gaya	Gurua	10	10	TractorPlo FourWheelTracRotav
ID_I3IPXS4DB7NE	Gaya	Gurua	15	15	TractorPlo FourWheelTracRotav
ID_4T8YQWXWHB4A	Nalanda	Noorsarai	60	60	TractorPlo WetTillagePudc
ID_W5MM9H353RL9	Vaishali	Garoul	10	5	TractorPlo
ID_6O44Z25H1JAV	Jamui	Khaira	12	12	TractorPlo
ID_VRI9LEL2W3DR	Nalanda	Rajgir	80	80	FourWheelTracRotav
ID_6YA9Y09O55LE	Jamui	Khaira	25	25	TractorPlo
ID_EDA8RK1CP60K	Nalanda	Noorsarai	20	10	WetTillagePudc

ID_920QSAHCN51N	Jamui	Khaira	25	14	TractorPlo
------------------------	-------	--------	----	----	------------

ID_LPERFIRDG4R1	Nalanda	Noorsarai	30	30	TractorPlo
------------------------	---------	-----------	----	----	------------

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 3870 entries, ID_GTFAC7PEVWQ9 to ID_KEPOQDTCZC6S
```

```
Data columns (total 43 columns):
```

#	Column	Non-Null Count	Dtype
0	District	3870 non-null	object
1	Block	3870 non-null	object
2	CultLand	3870 non-null	int64
3	CropCultLand	3870 non-null	int64
4	LandPreparationMethod	3870 non-null	object
5	CropTillageDate	3870 non-null	object
6	CropTillageDepth	3870 non-null	int64
7	CropEstMethod	3870 non-null	object
8	RcNursEstDate	3787 non-null	object
9	SeedingSowingTransplanting	3870 non-null	object
10	SeedlingsPerPit	3581 non-null	float64
11	NursDetFactor	3581 non-null	object
12	TransDetFactor	3581 non-null	object
13	TransplantingIrrigationHours	3677 non-null	float64
14	TransplantingIrrigationSource	3755 non-null	object
15	TransplantingIrrigationPowerSource	3367 non-null	object
16	TransIrriCost	2988 non-null	float64
17	StandingWater	3632 non-null	float64
18	OrgFertilizers	2535 non-null	object
19	Ganaura	1453 non-null	float64
20	CropOrgFYM	1196 non-null	float64
21	PCropSolidOrgFertAppMethod	2533 non-null	object
22	NoFertilizerAppln	3870 non-null	int64
23	CropbasalFerts	3682 non-null	object
24	BasalDAP	3327 non-null	float64
25	BasalUrea	2166 non-null	float64
26	MineralFertAppMethod	3870 non-null	object
27	FirstTopDressFert	3385 non-null	object
28	1tdUrea	3314 non-null	float64
29	1appDaysUrea	3314 non-null	float64
30	2tdUrea	1176 non-null	float64
31	2appDaysUrea	1170 non-null	float64
32	MineralFertAppMethod.1	3389 non-null	object
33	Harv_method	3870 non-null	object
34	Harv_date	3870 non-null	object
35	Harv_hand_rent	3618 non-null	float64
36	Threshing_date	3870 non-null	object
37	Threshing_method	3870 non-null	object
38	Residue_length	3870 non-null	int64
39	Residue_perc	3870 non-null	int64
40	Stubble_use	3870 non-null	object
41	Acre	3870 non-null	float64
42	Yield	3870 non-null	int64

```
dtypes: float64(14), int64(7), object(22)
```

```
memory usage: 1.3+ MB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	CultLand	CropCultLand	CropTillageDepth	SeedlingsPerPit	TransplantingIrr
count	3870.000000	3870.000000	3870.000000	3581.000000	
mean	28.527907	24.727132	4.488372	2.706507	
std	30.454218	27.994802	1.133044	7.624397	
min	1.000000	1.000000	1.000000	1.000000	
25%	12.000000	10.000000	4.000000	2.000000	
50%	20.000000	20.000000	4.000000	2.000000	
75%	35.000000	30.000000	5.000000	3.000000	
max	800.000000	800.000000	8.000000	442.000000	

Get the descriptive statistics

```
In [6]: desc = df.describe()
# Calculate the 95th percentile for numeric columns only
percentile_95 = df.select_dtypes(include=[float, int]).quantile(0.99)
# Append the 95th percentile to the descriptive statistics
desc.loc['99%'] = percentile_95
desc
```

```
Out[6]:
```

	CultLand	CropCultLand	CropTillageDepth	SeedlingsPerPit	TransplantingIrr
count	3870.000000	3870.000000	3870.000000	3581.000000	
mean	28.527907	24.727132	4.488372	2.706507	
std	30.454218	27.994802	1.133044	7.624397	
min	1.000000	1.000000	1.000000	1.000000	
25%	12.000000	10.000000	4.000000	2.000000	
50%	20.000000	20.000000	4.000000	2.000000	
75%	35.000000	30.000000	5.000000	3.000000	
max	800.000000	800.000000	8.000000	442.000000	
99%	120.000000	100.000000	8.000000	12.000000	

We can see there are outliers for most numeric features.

```
In [7]: # Check for duplicate rows  
df.duplicated().sum()
```

Out[7]: 0

```
In [8]: # Count missing values in each column  
df.isnull().sum()
```

```
Out[8]: District      0
        Block         0
        CultLand      0
        CropCultLand   0
        LandPreparationMethod 0
        CropTillageDate 0
        CropTillageDepth 0
        CropEstMethod  0
        RcNursEstDate  83
        SeedingSowingTransplanting 0
        SeedlingsPerPit 289
        NursDetFactor  289
        TransDetFactor 289
        TransplantingIrrigationHours 193
        TransplantingIrrigationSource 115
        TransplantingIrrigationPowerSource 503
        TransIrriCost  882
        StandingWater  238
        OrgFertilizers 1335
        Ganaura         2417
        CropOrgFYM       2674
        PCropSolidOrgFertAppMethod 1337
        NoFertilizerAppln 0
        CropbasalFerts   188
        BasalDAP          543
        BasalUrea        1704
        MineralFertAppMethod 0
        FirstTopDressFert 485
        1tdUrea          556
        1appDaysUrea     556
        2tdUrea          2694
        2appDaysUrea     2700
        MineralFertAppMethod.1 481
        Harv_method      0
        Harv_date        0
        Harv_hand_rent   252
        Threshing_date   0
        Threshing_method 0
        Residue_length   0
        Residue_perc     0
        Stubble_use      0
        Acre             0
        Yield            0
        dtype: int64
```

```
In [9]: # Calculate the percentage of missing values in each column
        (df.isnull().sum() * 100 / len(df)).round(2)
```

```

Out[9]: District      0.00
        Block         0.00
        CultLand      0.00
        CropCultLand  0.00
        LandPreparationMethod 0.00
        CropTillageDate 0.00
        CropTillageDepth 0.00
        CropEstMethod  0.00
        RcNursEstDate  2.14
        SeedingSowingTransplanting 0.00
        SeedlingsPerPit 7.47
        NursDetFactor  7.47
        TransDetFactor 7.47
        TransplantingIrrigationHours 4.99
        TransplantingIrrigationSource 2.97
        TransplantingIrrigationPowerSource 13.00
        TransIrriCost  22.79
        StandingWater  6.15
        OrgFertilizers 34.50
        Ganaura         62.45
        CropOrgFYM      69.10
        PCropSolidOrgFertAppMethod 34.55
        NoFertilizerAppln 0.00
        CropbasalFerts  4.86
        BasalDAP        14.03
        BasalUrea       44.03
        MineralFertAppMethod 0.00
        FirstTopDressFert 12.53
        1tdUrea         14.37
        1appDaysUrea    14.37
        2tdUrea         69.61
        2appDaysUrea    69.77
        MineralFertAppMethod.1 12.43
        Harv_method     0.00
        Harv_date       0.00
        Harv_hand_rent  6.51
        Threshing_date  0.00
        Threshing_method 0.00
        Residue_length  0.00
        Residue_perc    0.00
        Stubble_use     0.00
        Acre            0.00
        Yield           0.00
        dtype: float64

```

```
In [10]: df['NoFertilizerAppln'].value_counts()
```



```
Out[10]: NoFertilizerAppln  
2      2201  
3      1181  
1       481  
4         7  
Name: count, dtype: int64
```


```
In [11]: df['Harv_method'].value_counts()
```

```
Out[11]: Harv_method  
hand      3642  
machine    228  
Name: count, dtype: int64
```

Visualize Numerical Variables

Correlation Heatmap for Numeric Columns

```
In [12]: # Select and normalize only the numeric columns  
numeric_data = df.select_dtypes(include=['float64', 'int64'])  
scaler = MinMaxScaler()  
normalized_data = scaler.fit_transform(numeric_data)  
normalized_df = pd.DataFrame(normalized_data, columns=numeric_data.columns)  
  
# Compute the correlation matrix on normalized data  
correlation_matrix = normalized_df.corr()  
  
# Plot the correlation heatmap  
plt.figure(figsize=(16, 12))  
sns.heatmap(correlation_matrix, annot=True, cmap='viridis', square=True)  
plt.title("Correlation Heatmap for Normalized Numeric Features")  
plt.show()
```



There are no significant correlations between features except 'CultLand' and 'CropCultLand'.

```
In [13]: # Select numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

# Set up rows of 3 plots each
num_plots = len(numerical_columns)
num_rows = (num_plots // 3) + (num_plots % 3 > 0)

fig, axes = plt.subplots(num_rows, 3, figsize=(18, num_rows * 4))
axes = axes.flatten()

# Create histograms for each numerical column
for i, col in enumerate(numerical_columns):
```

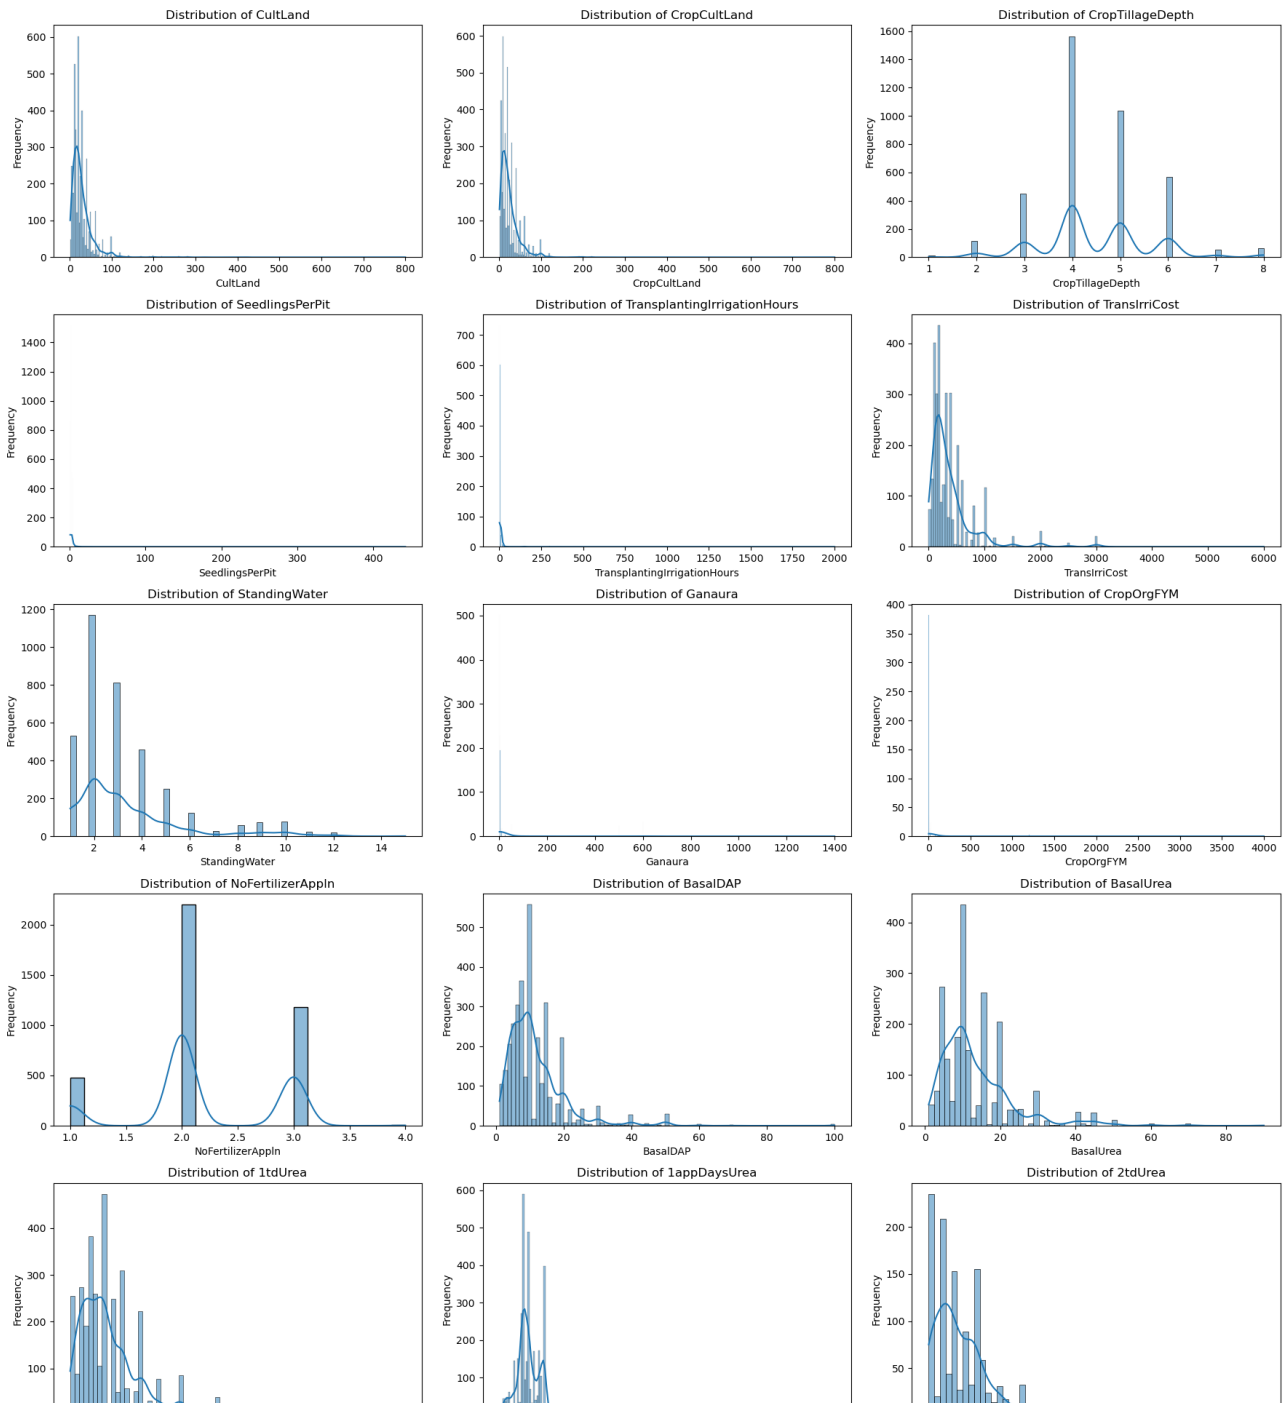
```

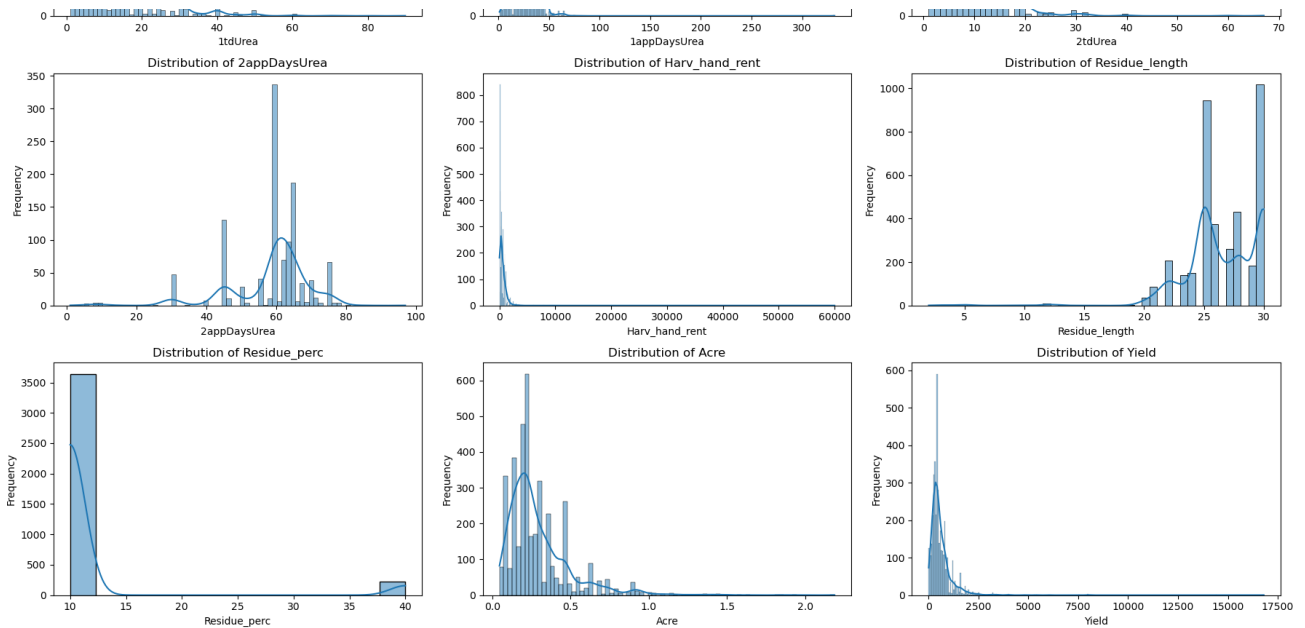
sns.histplot(df[col].dropna(), kde=True, ax=axes[i])
axes[i].set_title(f'Distribution of {col}')
axes[i].set_xlabel(col)
axes[i].set_ylabel('Frequency')

# Hide any extra subplots
for j in range(i + 1, num_rows * 3):
    axes[j].set_visible(False)

plt.tight_layout()
plt.show()

```





In [14]: numerical_columns

```
Out[14]: Index(['CultLand', 'CropCultLand', 'CropTillageDepth', 'SeedlingsPerPit',
               'TransplantingIrrigationHours', 'TransIrriCost', 'StandingWater',
               'Ganaura', 'CropOrgFYM', 'NoFertilizerAppln', 'BasalDAP', 'BasalUre
               a',
               '1tdUrea', '1appDaysUrea', '2tdUrea', '2appDaysUrea', 'Harv_hand_ren
               t',
               'Residue_length', 'Residue_perc', 'Acre', 'Yield'],
              dtype='object')
```

Let's look closer at some of the features

```
In [15]: # Calculate the 95th percentile for the 'Ganaura' column
q90 = df['Ganaura'].quantile(0.95)
print("95th Percentile (Upper Limit):", q90)
```

95th Percentile (Upper Limit): 80.0

```
In [16]: # Ganaura, CropOrgFYM
df['CropOrgFYM'].quantile(0.95)
```

Out[16]: 22.75

```
In [17]: # Calculate the 95th percentile
q95 = df['1appDaysUrea'].quantile(0.95)
print("95th Percentile (Upper Limit):", q95)
```

95th Percentile (Upper Limit): 45.0

```
In [18]: q95 = df['TransplantingIrrigationHours'].quantile(0.95)
```

```
print("95th Percentile (Upper Limit):", q95)
```

95th Percentile (Upper Limit): 15.0

```
In [19]: # Checking correlation between the two columns
correlation = df[['CropCultLand', 'Acre']].corr()
print(correlation)
```

	CropCultLand	Acre
CropCultLand	1.00000	0.39407
Acre	0.39407	1.00000

```
In [20]: # Checking correlation between the two columns
correlation = df[['CultLand', 'Acre']].corr()
print(correlation)
```

	CultLand	Acre
CultLand	1.000000	0.409604
Acre	0.409604	1.000000

```
In [21]: df['Yield'].quantile(0.90)
```

Out[21]: 1200.0

```
In [22]: df['Harv_hand_rent'].value_counts().sort_index()
```

```
Out[22]: Harv_hand_rent
1.0      1
2.0      2
3.0     11
4.0      4
5.0      7
..
6137.0    1
7221.0    1
7931.0    1
9300.0    2
60000.0    1
Name: count, Length: 131, dtype: int64
```

Let's inspect the number of fertilizer applications and their kind

```
In [23]: df['NoFertilizerAppln'].value_counts()
```

```
Out[23]: NoFertilizerAppln
2      2201
3      1181
1       481
4         7
Name: count, dtype: int64
```

```
In [24]: df[df['NoFertilizerAppln'] != 1][['FirstTopDressFert', '1tdUrea', '1appDaysUrea']]
```

```
Out[24]:
```

	FirstTopDressFert	1tdUrea	1appDaysUrea
ID			
ID_GTFAC7PEVWQ9	Urea	15.0	18.0
ID_TK40ARLSPOKS	Urea	20.0	39.0
ID_1FJY2CRIMLZZ	Urea	5.0	65.0
ID_I3IPXS4DB7NE	Urea	5.0	5.0
ID_4T8YQWXWHB4A	Urea	30.0	26.0
...
ID_DU6AHQ06QMXV	Urea	9.0	23.0
ID_PW2LN7ACB8MM	Urea DAP	2.0	30.0
ID_7ZZQ6R4XB4FK	Urea	12.0	45.0
ID_PVVDF6LK6FO8	Urea	6.0	45.0
ID_KEPOQDTCZC6S	Urea	10.0	28.0

3389 rows × 3 columns

```
In [25]: # Filter rows where 'FirstTopDressFert' does not contain "Urea" and 'NoFerti
filtered_df = df[(~df['FirstTopDressFert'].str.contains("Urea", na=False)) &

# Display the filtered result
print(filtered_df)
```

	FirstTopDressFert	1tdUrea	1appDaysUrea
ID			
ID_GRREKUJLG8N5	DAP	NaN	NaN
ID_RQ2X90R8F30U	DAP	NaN	NaN
ID_6VW9ED51TTXA	DAP	NaN	NaN
ID_6SNCEIE9GIKJ	DAP	NaN	NaN
ID_W53ULSZ3UZJK	DAP	NaN	NaN
...
ID_3KR27B00615L	DAP NPKS	NaN	NaN
ID_USJIQ3L9ZQLD	DAP	NaN	NaN
ID_9II4YBSXYK4Q	DAP	NaN	NaN
ID_2KJP9DE2ZBIY	DAP	NaN	NaN
ID_8I7QB5U74AUJ	DAP	NaN	NaN

[75 rows x 3 columns]

```
In [26]: df[df['NoFertilizerAppln']==3][['2tdUrea', '2appDaysUrea']]
```

Out [26]:

	2tdUrea	2appDaysUrea
ID		
ID_6O44Z25H1JAV	6.0	67.0
ID_6YA9Y09O55LE	7.0	58.0
ID_92OQSAHCN51N	12.0	65.0
ID_SJYVZSXJCX8S	1.0	60.0
ID_O99ZE30OJQ0E	1.0	60.0
...
ID_OQG31JUGU5JL	7.0	65.0
ID_RSC7O6MY665W	10.0	45.0
ID_DU6AHQ06QMXV	9.0	55.0
ID_PW2LN7ACB8MM	1.0	60.0
ID_PVVDF6LK6FO8	6.0	6.0

1181 rows × 2 columns

```
In [27]: df[(df['NoFertilizerAppln'].isin([3, 4])) & (df['2tdUrea'].isna() | df['2appDaysUrea'].isna())]
```

Out [27]:

	2tdUrea	2appDaysUrea
ID		
ID_X4024AVMS14D	6.0	NaN
ID_N1ELLJ1F32VU	NaN	NaN
ID_JI5NFWHR5G3H	4.0	NaN
ID_90IBO92HGZYD	5.0	NaN
ID_T5LYSK1D2TZX	NaN	NaN
ID_K9MMVQBHNSWY	NaN	NaN
ID_AXUCB1MQPA7Q	NaN	NaN
ID_Q5NRUPWTA5NX	3.0	NaN
ID_V7KHHGC42620	NaN	NaN
ID_4BUF8KA1KQN5	NaN	NaN
ID_S6XXFVWMWIEC	NaN	NaN
ID_FTCLUK815GGN	6.0	NaN
ID_LVFLM1AU48RZ	NaN	NaN
ID_6E1ZCZXS0M7H	6.0	NaN
ID_7W8Z7KFB21B8	NaN	NaN
ID_I35G50DPKNCY	NaN	NaN
ID_05LJ8YMX1TWL	NaN	NaN
ID_HZ9W937YJ3RI	NaN	NaN

In [28]: `df[(df['2tdUrea'].isna() | df['2appDaysUrea'].isna())][['2tdUrea', '2appDaysUrea']]`

Out [28]:

	2tdUrea	2appDaysUrea
ID		
ID_GTFAC7PEVWQ9	NaN	NaN
ID_TK40ARLSPOKS	NaN	NaN
ID_1FJY2CRIMLZZ	NaN	NaN
ID_I3IPXS4DB7NE	NaN	NaN
ID_4T8YQWXWHB4A	NaN	NaN
...
ID_NOASM3TXXY9	NaN	NaN
ID_7ZZQ6R4XB4FK	NaN	NaN
ID_RBYVUPRATVMW	NaN	NaN
ID_ARE9QWENJNJ2	NaN	NaN
ID_KEPOQDTCZC6S	NaN	NaN

2700 rows x 2 columns

Visualize Categorical Variables

```
In [29]: # Define the number of top categories to display (including "Other")
top_n = 10
categorical_columns = df.select_dtypes(include=['object']).columns

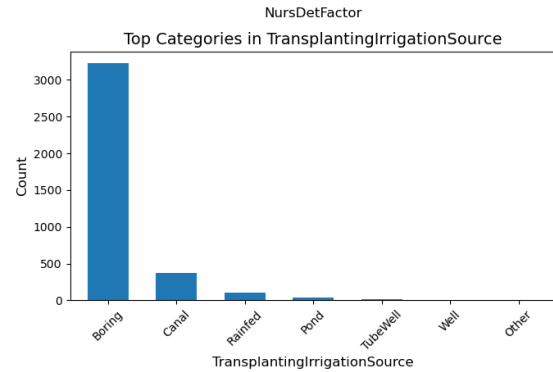

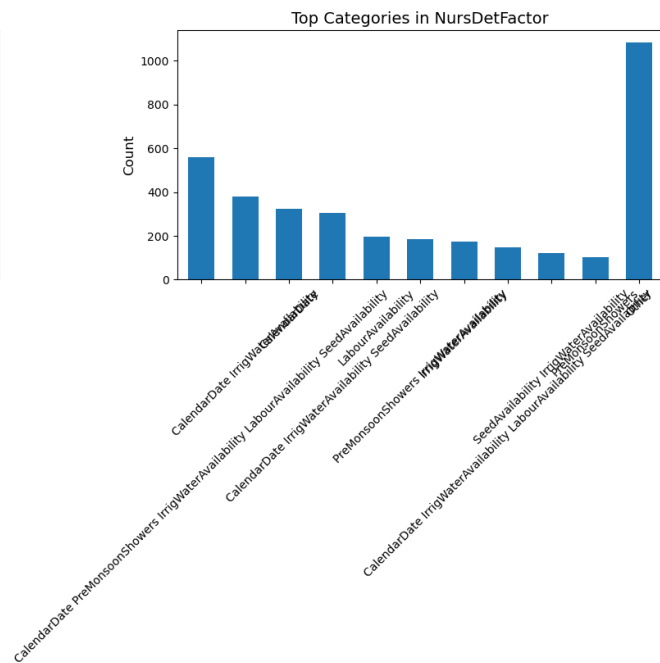


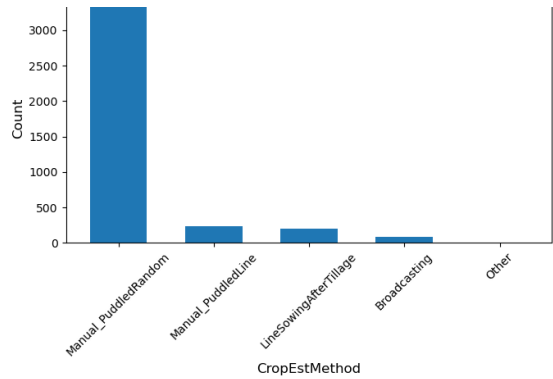
# Calculate rows and columns for the subplots grid
num_plots = len(categorical_columns)
num_rows = (num_plots // 2) + (num_plots % 2 > 0) # Two plots per row
fig, axes = plt.subplots(num_rows, 2, figsize=(16, num_rows * 8)) # Adjust
axes = axes.flatten()

# Plot top categories with an "Other" category as the 11th bar
for i, col in enumerate(categorical_columns):
    # Get the top N categories
    top_categories = df[col].value_counts().nlargest(top_n)



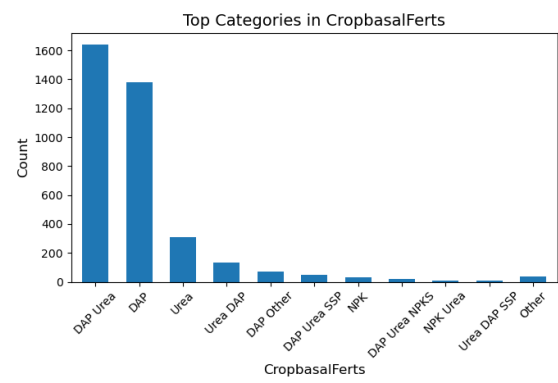
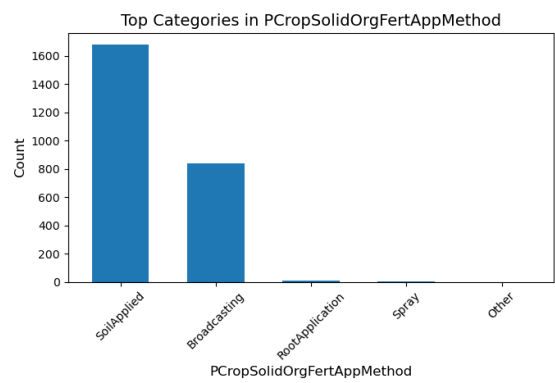
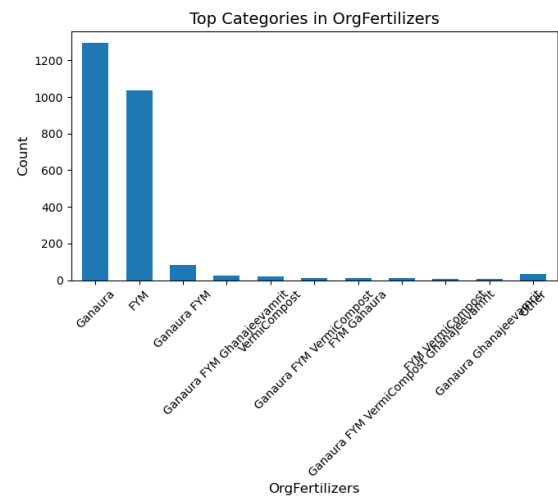
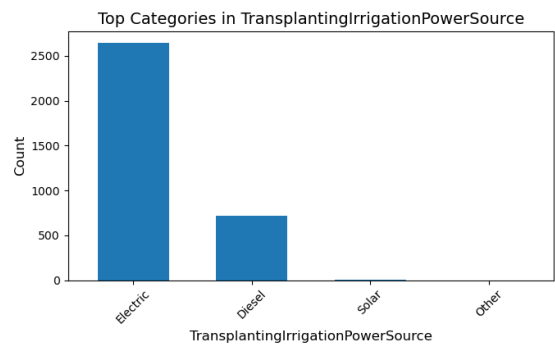
    # Calculate "Other" as the sum of all categories outside the top N
    other_count = df[col].value_counts()[top_n:].sum()
    top_categories['Other'] = other_count # Add "Other" category as the 11th



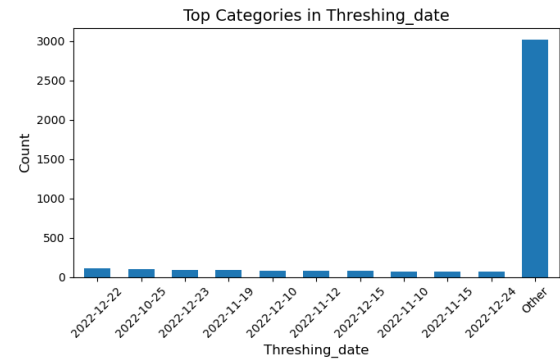
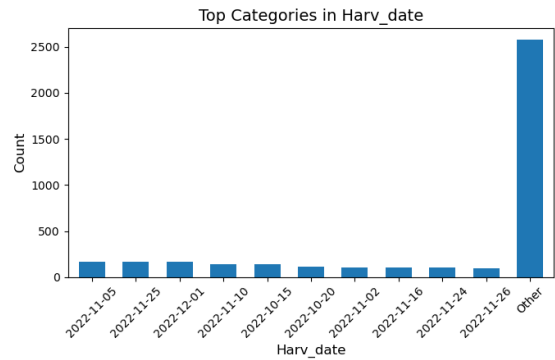
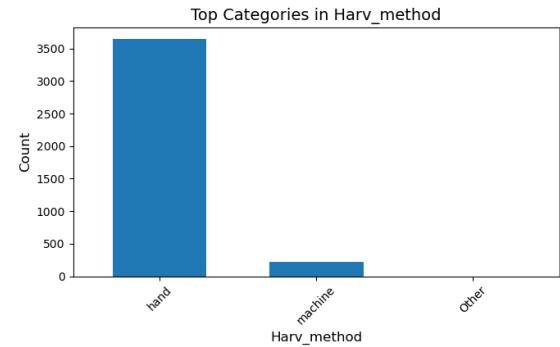

    # Plot
    top_categories.plot(kind='bar', ax=axes[i], width=0.6) # Adjust width t
```



Top Categories in RcNursEstDate



CalendarDate Seeding~
Seeding~
Calendar~
CalendarDate
TransDetFactor









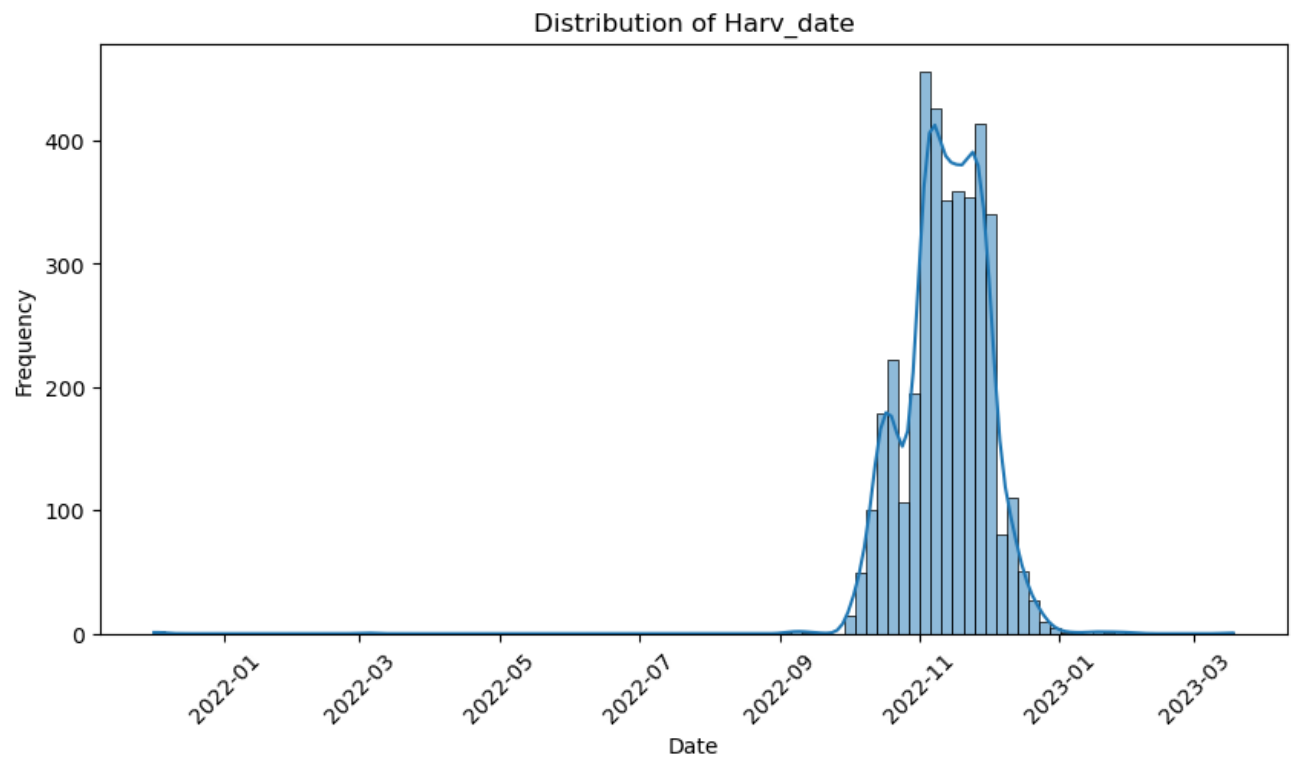

Visualize date columns


```
In [30]: # List of date columns
date_columns = ['CropTillageDate', 'RcNursEstDate', 'SeedingSowingTransplantDate']

# Temporarily convert date columns to datetime format
for col in date_columns:
    df[col] = pd.to_datetime(df[col])

# Plot the distribution of each date column
for col in date_columns:
    plt.figure(figsize=(10, 5))
    sns.histplot(df[col].dropna(), bins=100, kde=True) # Use KDE to see the distribution
    plt.title(f'Distribution of {col}')
    plt.xlabel('Date')
    plt.ylabel('Frequency')
    plt.xticks(rotation=45)
    plt.show()
```







```
In [31]: # Filter rows where 'Harv_date' is before October 1, 2022
df[df['Harv_date'] < '2022-09-01']
```

```
Out[31]:
```

	District	Block	CultLand	CropCultLand	LandPreparationMethod
ID					

ID_9P3DV08LL3SX	Jamui	Khaira	50	32	TractorPlough
-----------------	-------	--------	----	----	---------------

ID_YTZN9FE7PQUY	Nalanda	Rajgir	30	30	TractorPlough
-----------------	---------	--------	----	----	---------------

ID_RL2F5BMVBUAX	Jamui	Khaira	22	16	TractorPlough
-----------------	-------	--------	----	----	---------------

```
In [32]: df[df['Harv_date'] > '2023-01-01']
```

Out [32]:

	District	Block	CultLand	CropCultLand	LandPreparationMe
ID					
ID_PX8CNYP9YHPE	Vaishali	Chehrakala	8	6	TractorPl FourWheelTracRotav
ID_PATK559888IV	Gaya	Gurua	30	30	TractorPl FourWheelTracRotav
ID_ZAQF5TZH58PX	Nalanda	Rajgir	56	42	TractorPl FourWheelTracRotav
ID_J7PL485NOPNC	Nalanda	Rajgir	38	28	TractorPl FourWheelTracRotav
ID_X2XPIRO6ZYM	Nalanda	Rajgir	42	32	TractorPl FourWheelTracRotav
ID_WIXJZSB1JPAA	Gaya	Gurua	15	15	TractorPl FourWheelTracRotav
ID_RCQBDG SCHLXV	Jamui	Jamui	10	10	TractorPl
ID_Z0CHG2VQSJ31	Jamui	Jamui	20	20	TractorPl BullockPl
ID_2KJP9DE2ZBIY	Jamui	Jamui	18	5	WetTillagePud TractorPl BullockPl

```
In [33]: # Revert to original format if needed (optional)
for col in date_columns:
    df[col] = df[col].astype(str) # Convert back to string if needed
```

THE END

The rest of this notebook is failed model exploration.

```
In [ ]:
In [ ]:
In [ ]:
In [ ]:
```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

Preprocessing

Split data into training and testing

```
In [302... # Drop the index column
df = df.reset_index(drop=True)

X = df.drop('Yield', axis=1) # Features
y = df['Yield']              # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```